

Practical work 04 – 14/03/2023

Universal Representation Theorem

Model Selection

Objectives

The main objectives of this Practical Work for Week 4 are the following :

- a) Deepen your understanding of the Universal Representation Theorem.
- b) Play through an example of overfitting and determine the optimal model complexity by using hyper-parameter tuning based on 5-fold cross validation.

Submission

- **Deadline** : Tuesday 21 March, 3pm
- **Format** :
 - Exercise 1 (Universal Approximation Theorem) :
 - pdf with your calculation (handwritten) of the gradient of the MSE cost.
 - Jupyter notebook `function_approximation_stud.ipynb` completed with your solutions. The answers to the questions can be added either in the pdf report or in the notebook.
 - Exercise 2 (Model Selection)
 - Jupyter notebook.
 - Comments and results (plot with learning curve showing the results for different model complexities) either in the notebook or in a pdf-report.

Submission of all files in a single zip-file using the naming convention (for team of two students #1, #2) :

family name_given name #1- family name_given name #2.zip

Exercise 1 Function Approximation

In this exercise, you train a single (hidden) layer neural net to represent a given function $f : [0, 1] \rightarrow \mathbb{R}$. Since it is a regression problem, we will use the MSE cost function.

The MSE cost for a neural net with a 1d input x , a single hidden layer with n units and a linear output layer is given by

$$J_{\text{MSE}}(\theta) = \frac{1}{2m} \sum_{i=1}^m \left(y^{(i)} - \left(\sum_{k=1}^n w_{2,k} \sigma(w_{1,k} \cdot x^{(i)} + b_{1,k}) + b_2 \right) \right)^2 \quad (1)$$

The dataset is given by suitable x -values (in the interval $[0, 1]$) and associated function values $f(x)$, i.e. $\{(x^{(i)}, y^{(i)} = f(x^{(i)})) \mid i = 1, \dots, m\}$.

- Compute the formulas for gradient descent for this problem, i.e. compute the derivatives w.r.t. parameters $w_{1,k}, w_{2,k}, b_{1,k}, b_2$ and formulate the according update rules.
- Implement MBGD for this model in the notebook `function_approximation_stud.ipynb` by completing the code where indicated by

```
### START YOUR CODE ###
```

```
### END YOUR CODE ###
```

With the settings provided in the notebook (learning rate, batchsize, etc.) and the given dataset generated for the Beta-function (with $\alpha = \beta = 2.0$ and $m = 100$ samples) the learning should work quite well - see the learning curve (cost vs epochs) and the final MSE cost ($J_{\text{MSE}}(\theta_{\text{trained}}) \approx 7 \times 10^{-4}$).

- Now study the impact of different settings by looking at the learning curves (as a diagnostic tool) and the cost (obtained at the end of the last epoch) as performance measure. Consider :
 - Number of epochs : How many epochs are needed to see a reasonable fit ?
 - Learning rate : How large can you choose the learning rate ?
 - Number of neurons : How many neurons are needed for a sufficiently good approximation ?
 - Batch size : What happens when you increase the batch size ?

Can you improve the approximation as compared to the initial settings ?

- (Optional)** Now study different functions by generating new data with a different underlying function. Try e.g. the sine function with different frequencies. Start with a frequency $\omega = 1$ (ie. one cycle in the interval $[0, 1]$). Then proceed and investigate how large the (integer) frequency can be chosen to still obtain reasonable approximations. Possibly, also consider generating a larger dataset. Give an interpretation for why the learning breaks down at larger frequencies.

Exercise 2 Model Selection

The objective of this exercise is to build a classification systems to predict whether a student gets admitted into a University or not based on their results on two exams¹.

You have historical data from previous applicants that you can use as a training set. For each training example i , you have the applicant's scores on two exams $(x_1^{(i)}, x_2^{(i)})$ and the admissions decision $y^{(i)}$. Your task is to build a classification model that estimates an applicant's probability of admission based on the scores from those two exams.

In the notebook see `overfitting_stud.ipynb`, you'll find the code to load the data and further instructions.

- a) Complete the code where indicated by

```
### START YOUR CODE ###
```

```
### END YOUR CODE ###
```

Remarks : You can use the unit tests at the end of the notebook to test your implementation of the individual methods.

- b) Construct different (polynomial) models of different complexities (different degree i.e. parameter `order`). Train these models with the training set and determine the error rate on the training and the test set.
- c) Determine the model best suited for the problem at hand and justify why it is the best model.
- d) For all this use two different versions of the data :
- i) First version : `scores_train_1.csv` and `scores_test_1.csv` for training and testing, respectively.
 - ii) Second version : `scores_train_2.csv` and `scores_train_2.csv` for training and testing, respectively.

1. Data source : Andrew Ng - Machine Learning class Stanford