Practical work 06 – 25th of October 2022
Logistic Regression

Summary for the organisation :

— Submit the solutions of the practical work before Monday 12h00 next week in Moodle.
— **Rule 1**. Submit an archive (***.zip**!) with your Python notebooks (one per exercise), including datasets and all necessary files.
— **Rule 2**. The archive file name must contain the number of the practical work, followed by the family names of the team members by alphabetical order, for example `02_dupont_muller_smith.zip`. Put also the name of the team members in the body of the notebook (in first cell). Only one submission per team.
— **Rule 3**. We give a **fail** for submissions that do not compile (missing files are a common source of errors...). So, make sure that your whole notebooks give the expected solutions by clearing all cells and running them all before submitting.

Exercice 1   Classification to predict student admission

This is a continuation of the exercise of previous week, where the objective is to build a classification system to predict whether a student gets admitted into a university or not based on their results on two exams [1]. For each training example n, you have the applicants scores on two exams $(x_{n,1}, x_{n,2})$ and the admissions decision y_n. The dataset is illustrated on Figure 1.

a. Logistic regression classifier with linear decision boundary

The objective is to implement a classifier based on a logistic regression approach :

$$h_\theta(\mathbf{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

a) In a similar way as for the exercise of the previous week, read the training data from file `student-dataset-train.csv`. The first two columns are x_1 and x_2. The last column holds the class label y. Build the design matrix X as follow :

1. Data source : Andrew Ng - Machine Learning class Stanford

FIGURE 1 – Training data

$$X = \begin{pmatrix} 1 & x_{1,1} & x_{1,2} \\ 1 & \vdots & \vdots \\ 1 & x_{N,1} & x_{N,2} \end{pmatrix} \tag{1}$$

Check that the shape of X is $(100,3)$ and that the shape of y is $(100,)$.

b) Implement a z-norm normalization of the training set. You need to store the normalization values (μ, σ) for later as they will be needed to normalize the test set.

c) Implement a sigmoid function $g(z) = \frac{1}{1+e^{-z}}$. Use numpy to compute the exp so that your function can take numpy arrays as input. Check your implementation by plotting the sigmoid function.

d) Implement the hypothesis function $h_\theta(\mathbf{x})$. Hint : implement it so that the computation can take the full array X with $h(\mathbf{x})$ broadcasted to all training samples.

e) Implement the objective function $J(\theta)$:

$$J(\theta) = \frac{1}{N} \sum_{n=1}^{N} y_n \log h_\theta(\mathbf{x_n}) + (1 - y_n) \log(1 - h_\theta(\mathbf{x_n}))$$

f) In a similar way as in PW02 and PW03, implement the gradient ascent with the update rule :

$$\theta_i \leftarrow \theta_i + \alpha \frac{1}{N} \sum_{n=1}^{N} (y_n - h_\theta(\mathbf{x}_n)) x_{n,i}$$

g) Test your implementation by running a gradient ascent. Hints : use a small α. e.g. 0.001, store the evolution of the objective function $J(\theta)$ during the epochs to make a plot, use a large number of epochs, e.g. 2000000.

2

h) Compute the correct classification rate on `student-dataset-test.csv` after convergence as you have an estimator of the posterior probabilities with

$$P(y_n = 1|\mathbf{x_n}; \theta) = h_\theta(\mathbf{x_n})$$
$$P(y_n = 0|\mathbf{x_n}; \theta) = 1 - h_\theta(\mathbf{x_n})$$

This means that you can take the decisions $\hat{y}_n = 1$ if $h_\theta(\mathbf{x_n}) \geq 0.5$ and $\hat{y}_n = 0$ if $h_\theta(\mathbf{x_n}) < 0.5$.

i) Draw the decision boundary of your system on top of the scatter plot of the testing data.

j) Compare the performance of the logistic regression system with the ones of previous's week.

b. Optional - Stochastic gradient ascent

Redo the experiments of 2.a with a stochastic gradient ascent.

c. Logistic regression classifier with non-linear decision boundary

Redo the experiments of 2.a by increasing the complexity of the model in order to have a non-linear decision boundary :

$$h_\theta(\mathbf{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1 x_2 + \ldots)$$

d. Using SciKit Learn

Redo one of the exercise a. or c. using SciKit Learn.

a) Read the documentation of the function `SGDClassifier()` available in the toolkit SciKit Learn. This function implements stochastic gradient descent training for different linear systems such as Logistic Regression. For a logistic regression, the `loss` parameter should be set to `'log'`.

b) Use the `fit()` and `predict()` methods of this classifier on the student data.

c) Compute the performances and compare it to your own implementations.

Exercice 2   Optional - Classification on MNIST

Using the SciKit Learn toolkit, train a multi-class logistic regression on the MNIST problem and compare the performances with the KNN of PW2.

For the "daring" ones, implement a multi-class version based on your own implementation and run it against MNIST data.

Exercice 3   Review questions

a) Why do we have a gradient ascent in the case of logistic regression while we had a gradient descent with linear regression? Can we convert the gradient ascent of logistic regression into a gradient descent? If yes, how?

b) Assuming a logistic regression with a linear decision boundary taking as input samples in two dimensions (x_1, x_2), in which case do we get 0.5 as output of the classification system? Express your answer with an equation.

c) What is the computational trick to avoid numerical problems in the computation of $J(\theta)$ for the logistic regression? In which situations (for what type of inputs) do we risk to observe such numerical problems?

d) A logistic regression can classify between 2 classes. How can we build a multi-class (with K classes) system with logistic regression?